



ECE408/CS483/CSE408

Applied Parallel Programming

Lecture 15

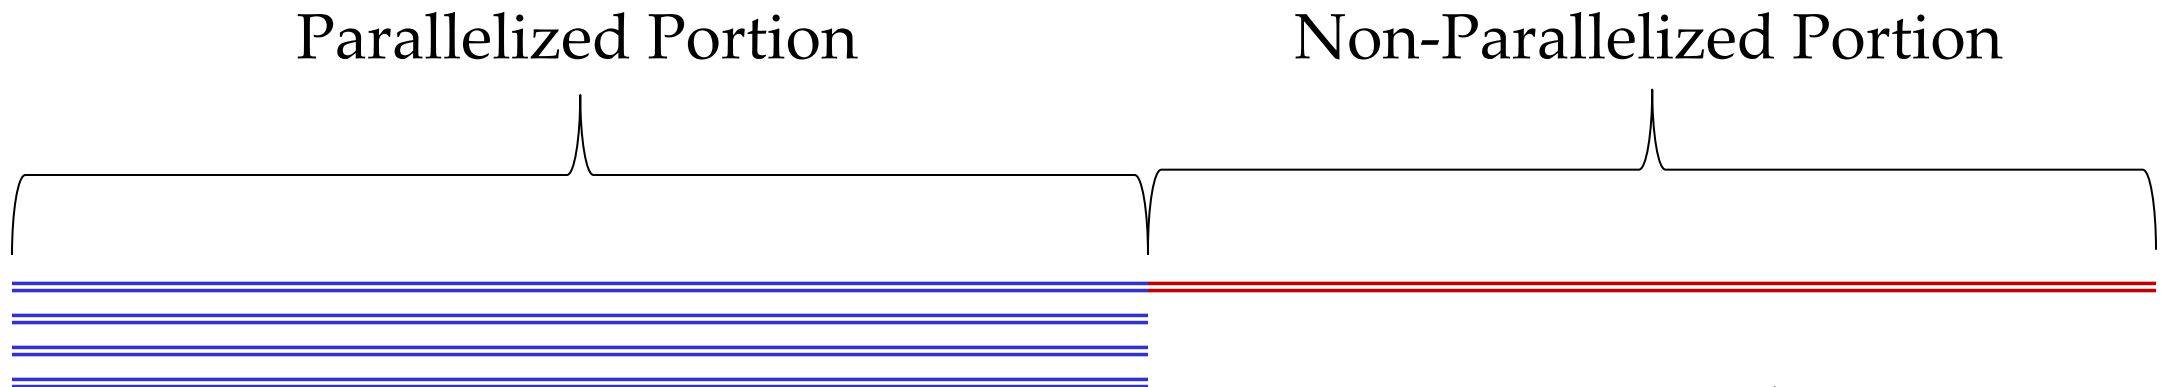
Parallel Computation Patterns – Reduction Trees

Course Reminders

- Exam 1: Tues Oct 10 @ 7pm
 - Room assignments are available on Canvas
 - 1 sheet of handwritten notes
- Project Milestone 1: CPU Convolution, profiling
 - Due tomorrow (Friday, Oct 13)
 - Project details are posted in GitHub
- Lab 5.1 due Oct 20, will be made available next week

Amdahl's Law

- First articulated by Gene Amdahl, computer pioneer, 1967



$$Speedup = \frac{1}{(1 - k) + (\frac{k}{p})}$$

k = parallelized portion

p = parallel resources

Assuming a fixed problem size

Reductions are everywhere...

often the final stage of parallel computation

- “Reduce” a set of input values into a singular value
 - Max
 - Min
 - Sum
 - Product
 - Dot Product
- Basically, any operator/function that satisfies the following:
 - Is associative and commutative
 - Has a well-defined identity value (e.g., 0 for sum)

Partition and Summarize

- A commonly used strategy for processing large input data sets
 - If the basic operation is associative and commutative (i.e., reorderable)
 - Partition the data set into smaller chunks
 - Have each thread to process a chunk
 - Use a reduction tree to summarize the results from each chunk into the final answer
- Big Data cloud frameworks (from Google and others) provide support for this pattern

Reduction enables other techniques


- Reduction is often needed as the final stage of parallel computation
- For example: Privatization
 - Multiple threads write into an output location
 - Replicate the output location so that each thread has a private output location
 - Use a reduction tree to combine the values of private locations into the original output location

An efficient sequential reduction algorithm performs $O(N)$ operations

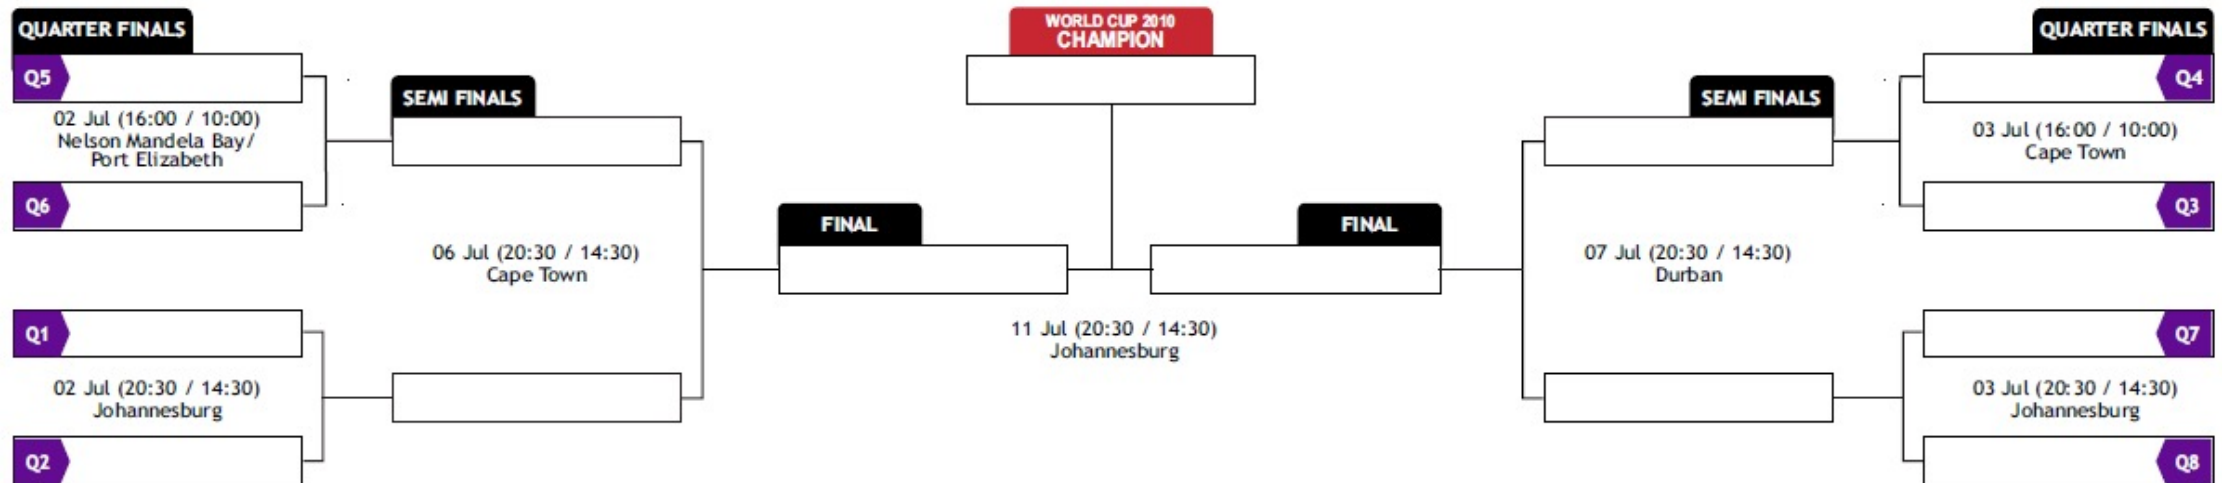
- Initialize the result as an identity value I for the reduction operation \square
 - Smallest possible value for max reduction
 - Largest possible value for min reduction
 - 0 for sum reduction
 - 1 for product reduction
- Iterate through the input and perform the reduction operation between the result value and the current input value

```
result  $\leftarrow$  I;  
for each value X in input  
    result  $\leftarrow$  result  $\square$  X;
```

A Parallel Reduction can be done in $\log(N)$ Steps



Tournaments Use Reduction with “max”



(A more artful rendition of the reduction tree.)

The Parallel Algorithm is Work Efficient

For N input values, the number of operations is

$$\frac{1}{2}N + \frac{1}{4}N + \frac{1}{8}N + \cdots + \frac{1}{N}N = \left(1 - \frac{1}{N}\right)N = N - 1.$$

The parallel algorithm shown is work-efficient: requires the same amount of work as a sequential algorithm (overheads might be different)

But requires a lot of resources...

For N input values, the number of steps is $\log(N)$.

With enough execution resources,

- $N=1,000,000$ takes 20 steps!


How much parallelism do we need?

- On average, $(N-1)/\log(N)$.
50,000 in our example.
- But peak is $N/2$!
500,000 in our example.


A Sum Reduction Example

- Parallel implementation:
 - Sum two values in each thread in each step
 - Recursively halve the # of threads
 - Takes $\log(n)$ steps for n elements, requires $n/2$ threads
- Assume an in-place reduction using shared memory
 - The original vector is in device global memory
 - The shared memory is used to hold a partial sum vector
 - Each step brings the partial sum vector closer to the sum
 - The final sum will be in element 0
 - Global memory traffic should be minimal due to use of shared memory

Initial Data Mapping for a Reduction



A Sum Example (Values Instead of Indices)




The Reduction Steps

```
// Stride is distance to the next value being
// accumulated into the threads mapped position
// in the partialSum[] array
for (unsigned int stride = 1;
     stride <= blockDim.x;  stride *= 2)
{
    __syncthreads();
    if (t % stride == 0)
        partialSum[2*t] += partialSum[2*t+stride];
}
```

Why do we need `__syncthreads()`?

Example Without __syncthreads



Several Options after Blocks are Done

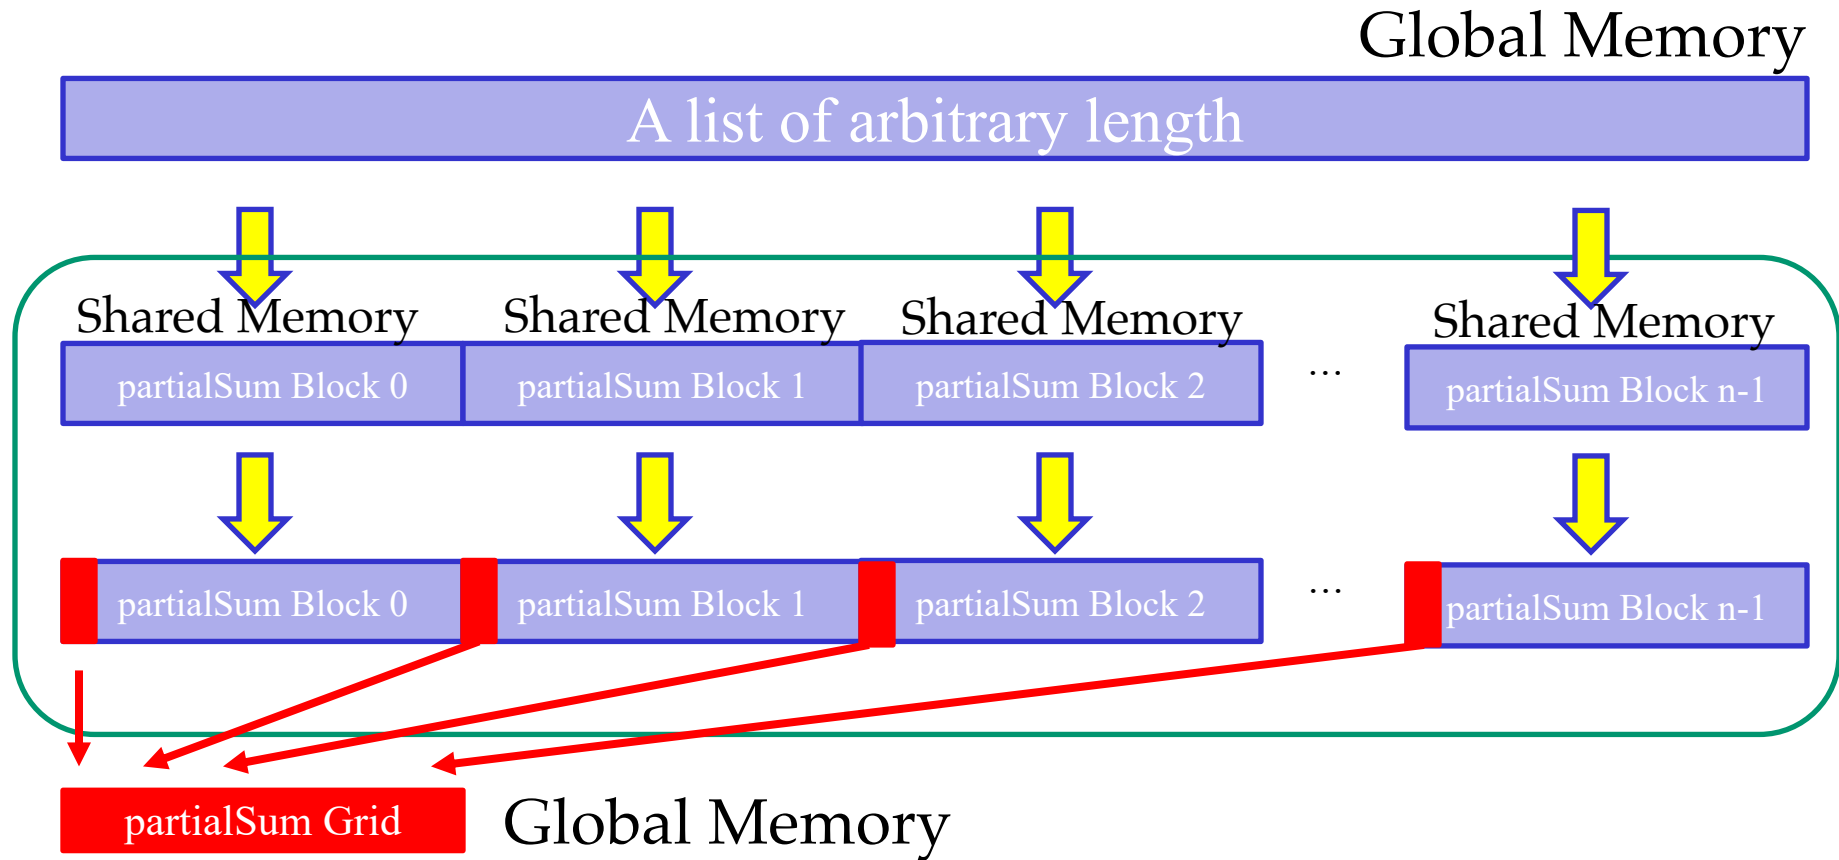
After all reduction steps, **thread 0**

- **writes** block's **sum** from **partialSum[0]**
- **into global vector** indexed by **blockIdx.x**.

Vector has **length $N / (2 * \text{numBlocks})$** .

- If small, **transfer** vector **to host** and **sum** it up **on CPU**.
- If large, **launch kernel again** (and again).
(Kernel can also accumulate to a global sum using atomic operations, to be covered soon.)

“Segmented Reduction”



Copy back to host and host to finish the work.

Analysis of Execution Resources

All threads active in the **first step**.

In all **subsequent steps**, two control flow paths:

- **perform addition, or do nothing.**
- Doing nothing still consumes execution resources.

At most half of threads perform addition after first step

- (all threads with odd indices disabled after first step).
- **After fifth step**, entire warps do nothing:
poor resource utilization, but no divergence.
- **Active warps have** only **one active thread**.

Up to five more steps (if limited to 1024 threads).

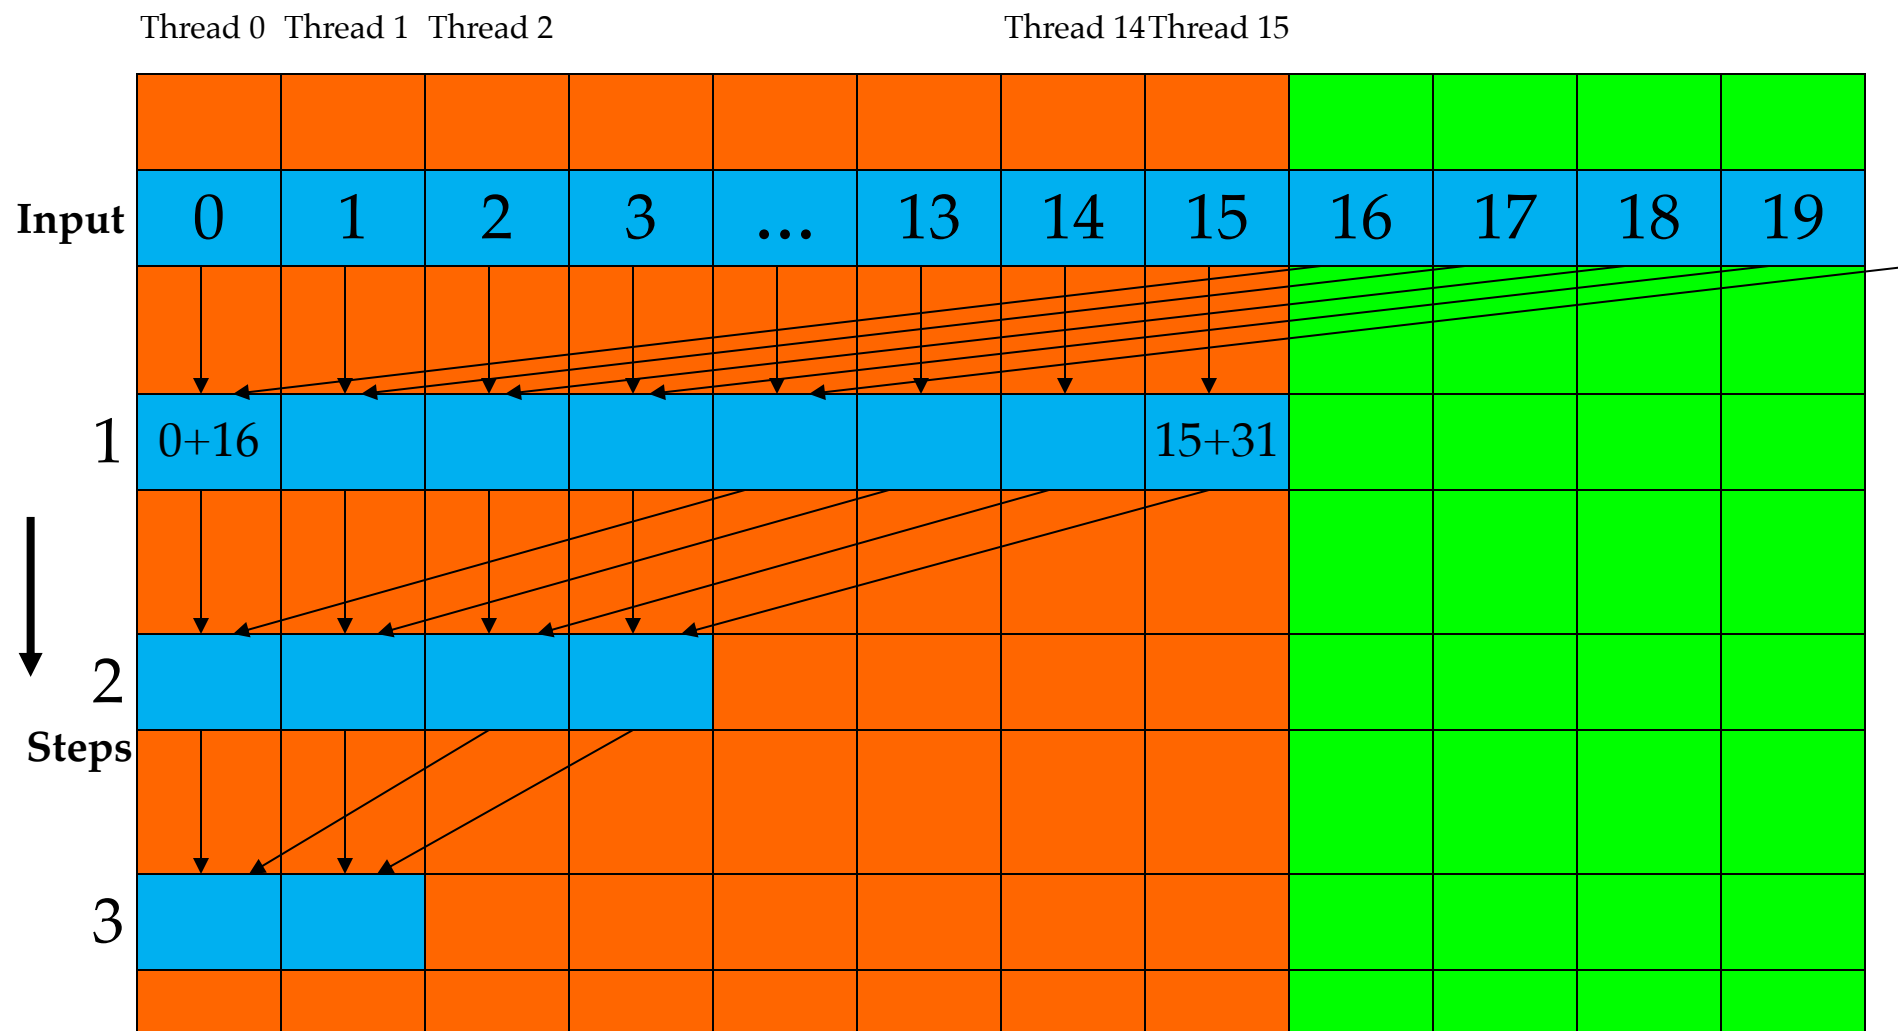
A Better Strategy

Let's try this approach:

- **in each step,**
- **compact** the partial sums
- **into** the **first locations**
- in the **partialSum** array

Doing so **keeps** the **active threads consecutive**.

Illustration with 16 Threads



A Better Reduction Kernel

```
for (unsigned int stride = blockDim.x; stride >= 1; stride /= 2)
{
    __syncthreads();
    if (t < stride)
        partialSum[t] += partialSum[t+stride];
}
```

Again: Analysis of Execution Resources

Given 1024 threads,

- Block loads 2048 elements to shared memory.
- **No branch divergence** in the **first six steps**:
 - 1024, 512, 256, 128, 64, and 32 consecutive threads active;
 - threads in each warp either all active or all inactive
- **Last six steps** have **one active warp** (branch divergence for last five steps).


Parallel Algorithm Overhead

```
__shared__ float partialSum[2*BLOCK_SIZE];

unsigned int t = threadIdx.x;


unsigned int start = 2*blockIdx.x*blockDim.x;
partialSum[t] = input[start + t];
partialSum[blockDim+t] = input[start+ blockDim.x+t];
for (unsigned int stride = blockDim.x; stride >= 1; stride /= 2)
{
    __syncthreads();
    if (t < stride)
        partialSum[t] += partialSum[t+stride];
}
```


Parallel Execution Overhead



Although the number of “operations” is N , each “operation” involves much more complex address calculation and intermediate result manipulation.

If the parallel code is executed on a single-thread hardware, it would be slower than the code based on the original sequential algorithm.



Further Improvements

The **problem has low compute-intensity**

- **one operation** for every **4B** value **read**;
- so **focus on memory coalescing and**
avoiding poor **computational resource use**.
- Experiment with Hyper-parameters:
 - Threads per block
 - Blocks per SM
 - Thread granularity (3-to-1 reduction instead of 2-to-1 reduction)

Dealing with Narrowing Parallelism

Smaller blocks might seem attractive:

- when one warp is active,
- each SM has one warp per block.

But there are probably better ways. For example,

- **stop reducing at 32 elements** (or at 64, or 128), and
- hand off to the **next kernel**.

Work Until the Data is Exhausted!

Say there are 8 SMs, so 16 blocks.

1. **Divide** the whole **dataset into 16 chunks**.
2. **Read** enough **to fill shared memory**.
3. **Compute** ... only **until** some **threads not needed**.
4. Then **load more data!**
5. **Repeat** until the data are exhausted,
6. THEN let parallelism drop.
(Gather 16 values on host and reduce them.)

Two vertical lines, one blue and one orange, are positioned on the left side of the slide.

ANY MORE QUESTIONS?
READ CHAPTER 5